



# GPU Acceleration with the C++ Standard Library

# Intended Audience

**Audience:** students, developers, researchers and practitioners interested in developing portable HPC applications using ISO C++

**Prerequisites:** experience with C++11 (lambdas, STL algorithms)

# GPU Acceleration with the C++ Standard Library

C++ Prerequisites

Fundamentals of ISO C++ Parallelism

Indexing, Ranges & Views

Interactive Materials

# C++ Prerequisites

# ISO C++ lambdas

Lambdas simplify the creation of function objects. This...

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [s,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);
```

# ISO C++ lambdas

Lambdas simplify the creation of function objects. This...

...is equivalent to...

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [s,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

The [...] is called the "lambda capture" and controls how variables are stored within the lambda object:

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [s,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

The [...] is called the "lambda capture" and controls how variables are stored within the lambda object:

- [s] captures s "by value" (makes a copy)

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [s,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

The [...] is called the "lambda capture" and controls how variables are stored within the lambda object:

- [s] captures s "by value" (makes a copy)
- [&v] captures v "by reference" (stores a pointer)

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [s,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

Lambda captures support "capture defaults" that capture all variables used within the lambda:

- `[&,s]` captures `s` "by value" (makes a copy) and all other used variables "by reference" (store pointers)
- `[=,&v]` captures `v` "by reference" (stores a pointer to `v`) and all other used variables "by value" (copy them)

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [&,s](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

Lambda captures support "capture defaults" that capture all variables used within the lambda:

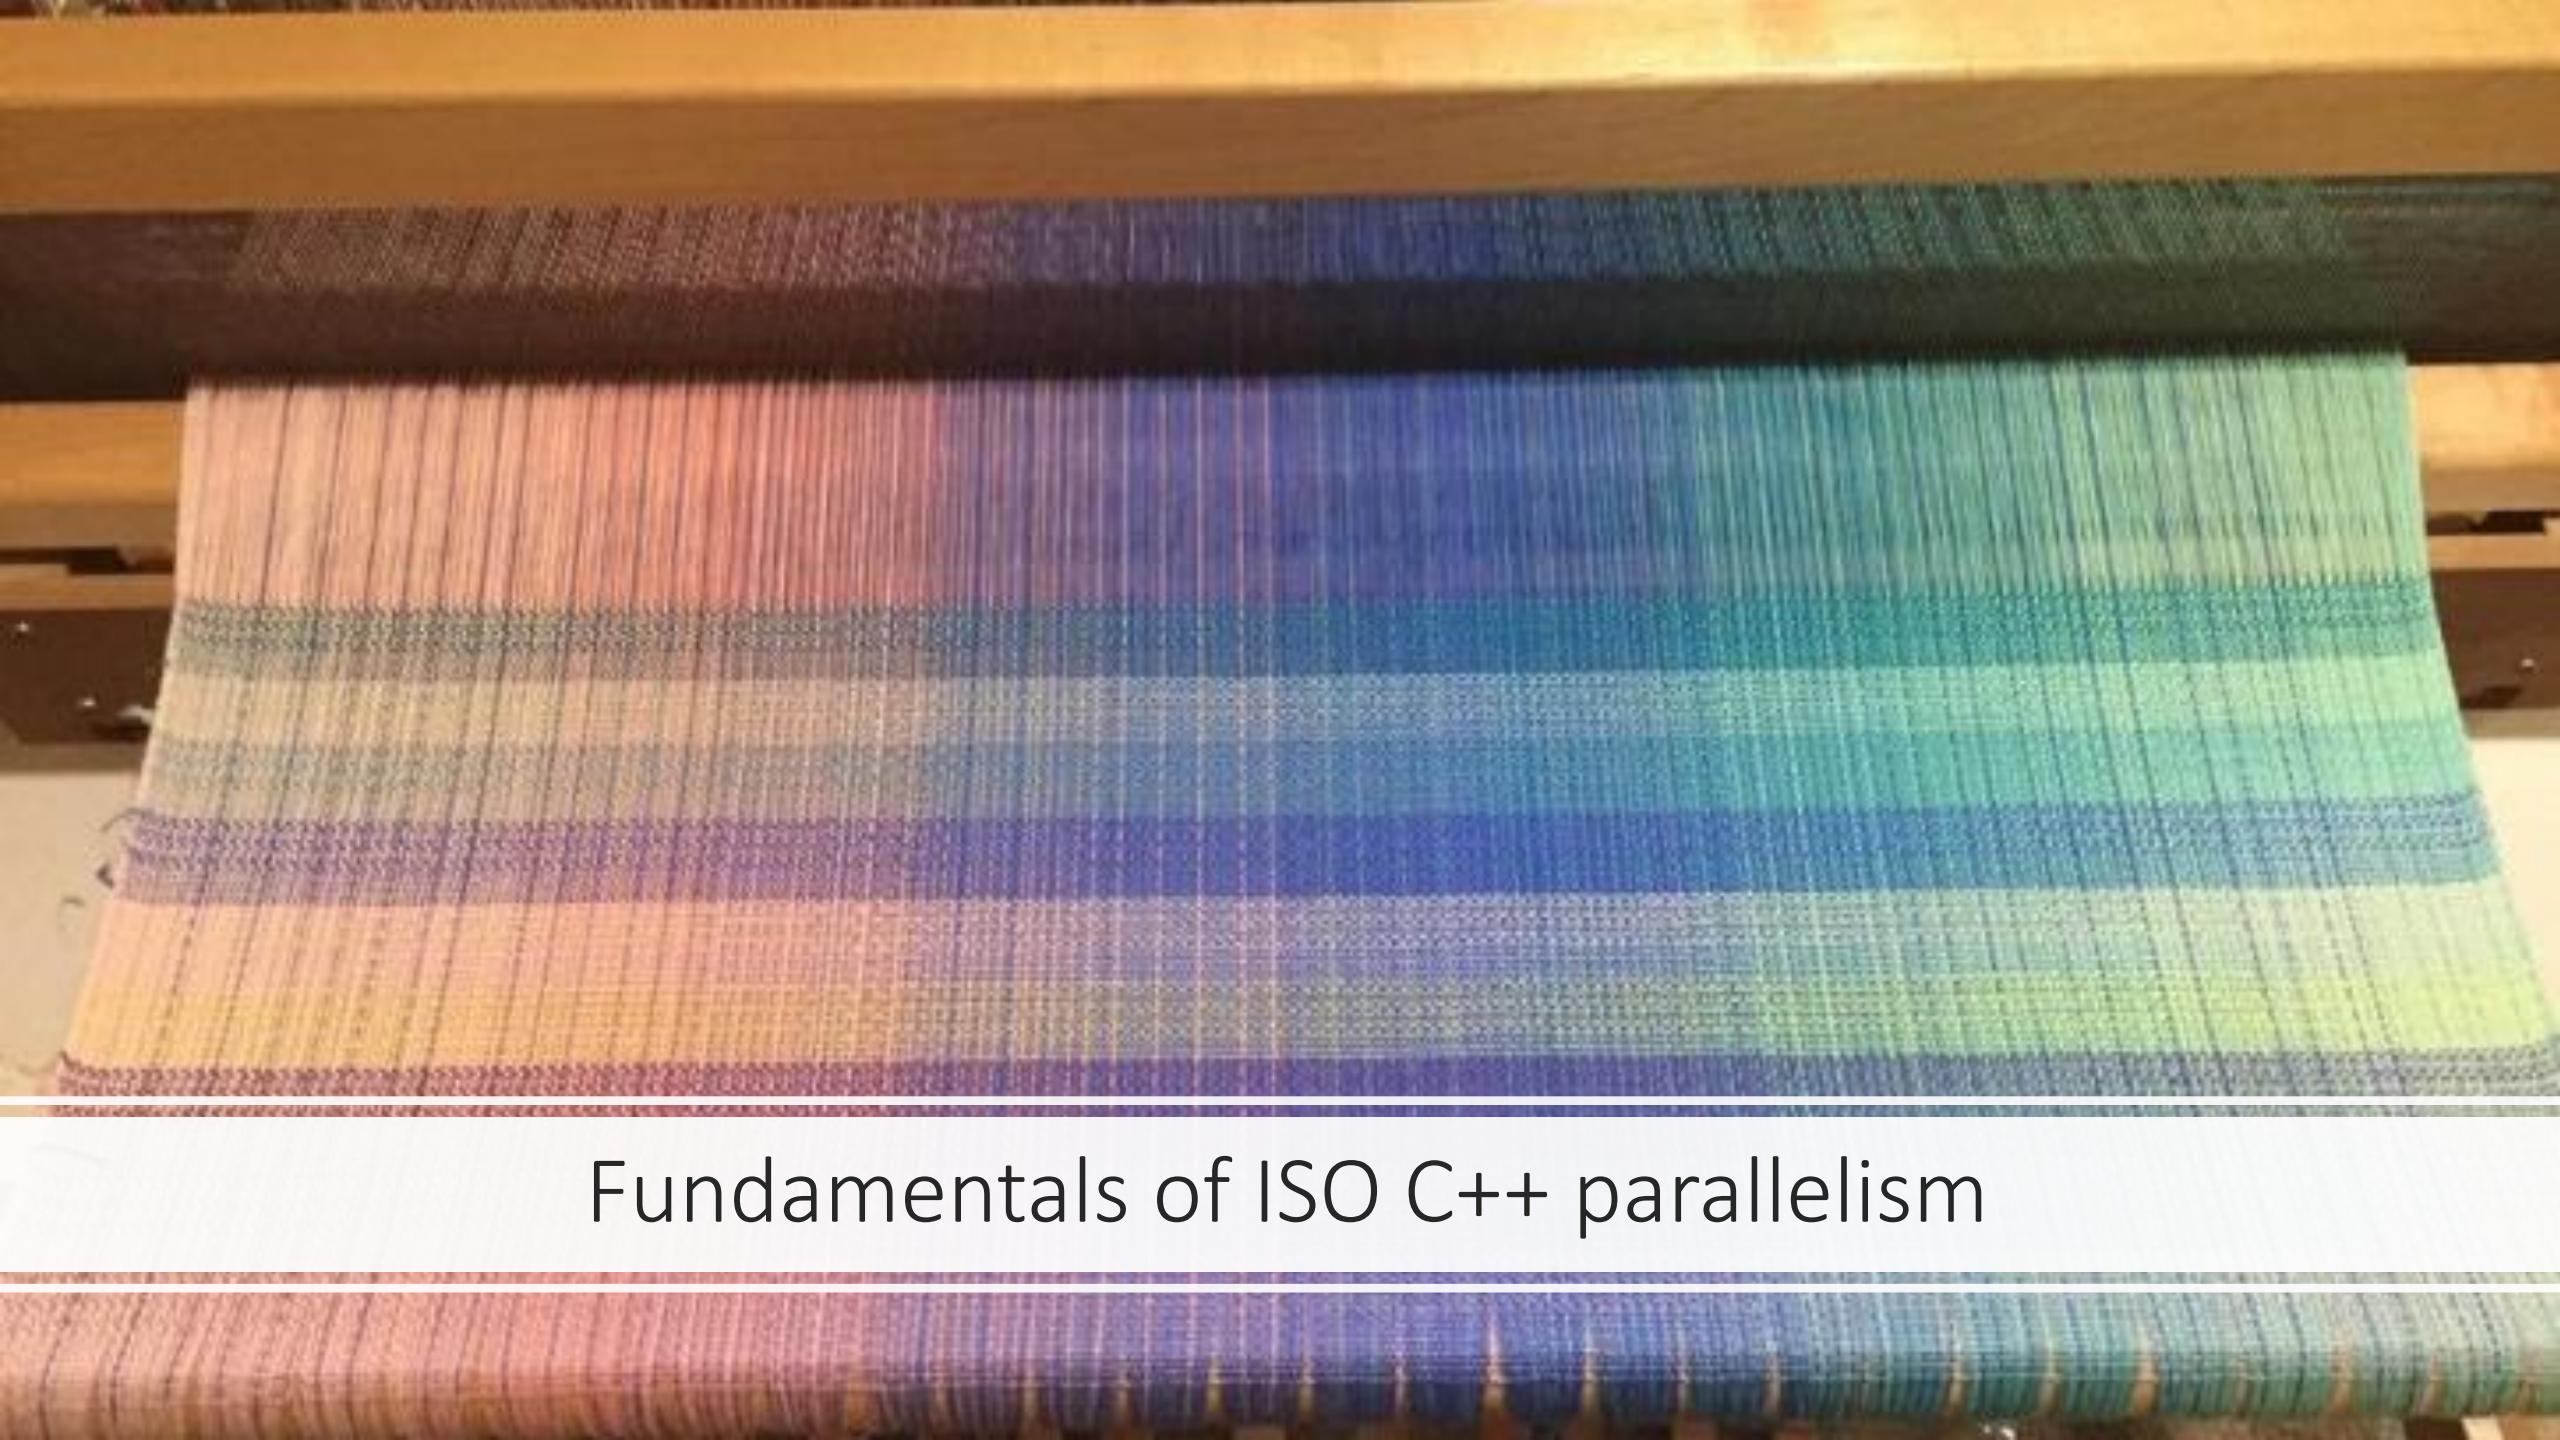
- `[&,s]` captures `s` "by value" (makes a copy) and all other used variables "by reference" (store pointers)
- `[=,&v]` captures `v` "by reference" (stores a pointer to `v`) and all other used variables "by value" (copy them)

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [=,&v](int idx) { return v[idx] * s; };  
assert(f(1) == 4);  
  
struct __unnamed {  
    double s;  
    std::vector<double>& v;  
    double operator()(int idx) {  
        return v[idx] * s;  
    }  
};  
__unnamed f{s, v};  
assert(f(1) == 4);
```

# ISO C++ lambdas

Lambda captures support creating and assigning to new variables for use within the lambda:

```
std::vector<double> v = {1, 2, 3, 4};  
double s = 2.;  
auto f = [a = s, x = v.data()](int idx) {  
    return x[idx] * a;  
};  
assert(f(1) == 4);
```



Fundamentals of ISO C++ parallelism

# ISO C++ algorithms

In C++, containers can be processed by **for** loops...

```
std::vector<double> v = {1, 2, 3, 4}, w(4);
for (int i = 0; i < 4 ; ++i) {
    w[i] = 2. * v[i];
}
```

# ISO C++ algorithms

In C++, containers can be processed by **for** loops...

```
std::vector<double> v = {1, 2, 3, 4}, w(4);
for (int i = 0; i < 4 ; ++i) {
    w[i] = 2. * v[i];
}
```

... or with standard template library (STL) algorithms, which are often more succinct.

```
std::transform(begin(v), end(v), begin(w),
              [](&const double& el) {
                  return 2. * el;
});
```

# ISO C++ parallel algorithms

## Programming model introduced in C++17

```
std::vector<double> v = {1, 2, 3, 4}, w(4);
std::transform(std::execution::par, begin(v), end(v), begin(w),
    [](const double& el) { return 2. * el; });
```

# ISO C++ parallel algorithms

## Programming model introduced in C++17

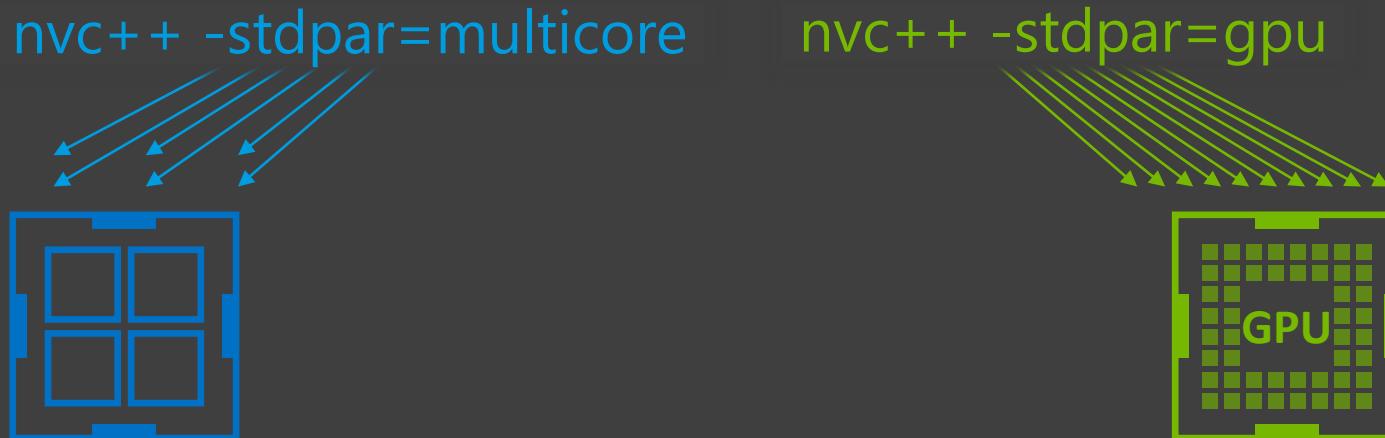
```
std::vector<double> v = {1, 2, 3, 4}, w(4);
std::transform(std::execution::par, begin(v), end(v), begin(w),
              [] (const double& el) { return 2. * el; });
```



# ISO C++ **parallel** algorithms

## Compiler selects target for parallel execution

```
std::vector<double> v = {1, 2, 3, 4}, w(4);
std::transform(std::execution::par, begin(v), end(v), begin(w),
              [](const double& el) { return 2. * el; });
```



# ISO C++ **parallel** algorithms Hybrid (CPU / GPU) program execution

```
std::vector<double> v = {1, 2, 3, 4}, w(4);

// Data is first processed sequentially on the host (CPU)
std::transform(begin(v), end(v), begin(w),
               [](const double& el) { return 2. * el; });

// Then, the same data is processed in parallel, e.g. on a GPU
std::transform(std::execution::par, begin(v), end(v), begin(w),
               [](const double& el) { return 2. * el; });
```

- Same data can be accessed from the CPU and from the GPU.
- **Memory transfer** is implicit.
- Use of a **unified (managed)** memory model.

# ISO C++ **parallel** algorithms

## Accelerator support limitation

Stack variable “a” and global variable “b” are captured by reference (**&**).

Accelerators read it remotely from the CPU thread stack.

- Non-coherent HW (PCIe):  
**not supported**
- Coherent HW (Grace Hopper): poor performance.
- Note: this is a problem for stack data and globals only, not for heap data (the vector).

```
double b;  
void multiply_with(vector<double>& v, double a) {  
    std::for_each(std::execution::par,  
                 begin(v), end(v),  
                 [&](double& x) { x *= a + b; })  
};  
}
```

# ISO C++ parallel algorithms

## Accelerator support limitation

Stack variable “a” is captured by reference (`&`). This is problematic (non-supported or slow).

Solution: Stack variable “a” is captured by value (`=`) and copied to the accelerator.

```
void multiply_with(vector<double>& v, double a) {  
    std::for_each(std::execution::par,  
                 begin(v), end(v),  
                 [&](double& x) { x *= a; })  
};  
}
```

```
void multiply_with(vector<double>& v, double a) {  
    std::for_each(std::execution::par,  
                 begin(v), end(v),  
                 [=](double& x) { x *= a; })  
};  
}
```

# ISO C++ parallel algorithms

## References

CppCon talks:

- Thomas Rodgers, *Bringing C++ 17 Parallel Algorithms to a standard library near you*, 2018
- Olivier Giroux, *Designing (New) C++ Hardware*, 2017
- Dietmar Kühl, *C++17 Parallel Algorithms*, 2017
- Bryce Adelstein Lelbach, *The C++17 Parallel Algorithms Library and Beyond*, 2016

GTC talks:

- Simon McIntosh-Smith et al., *How to Develop Performance Portable Codes using the Latest Parallel Programming Standards*, Spring 2022
- Jonas Latt, *Porting a Scientific Application to GPU Using C++ Standard Parallelism*, Fall 2021
- Jonas Latt, *Fluid Dynamics on GPUs with C++ Parallel Algorithms: State-of-the-Art Performance through a Hardware-Agnostic Approach*, Spring 2021

# C++ Parallel Algorithms in C++17 & C++20

See <https://en.cppreference.com/w/cpp/algorithm>

## Iteration & Transform

`std::for_each, std::for_each_n`  
`std::transform, std::transform_reduce`  
`std::transform_inclusive_scan, std::transform_exclusive_scan`

## Reductions

`std::reduce, std::transform_reduce`  
`std::exclusive_scan, std::inclusive_scan`  
`std::adjacent_difference`  
`std::all_of, std::any_of, std::none_of`  
`std::count, std::count_if`  
`std::is_sorted, std::is_sorted_until, std::is_partitioned`  
`std::is_heap, std::is_heap_until`  
`std::max_element, std::min_element, std::minmax_element`  
`std::equal, std::lexicographical_compare`

## Searching

`std::find, std::find_if, std::find_if_not, std::find_end, std::find_first_of`  
`std::adjacent_find, std::mismatch`  
`std::search, std::search_n`

## Memory movement & Initialization

`std::copy / copy_if / copy_n / move / uninitialized_...`  
`std::fill, std::fill_n, std::uninitialized_...`  
`std::generate, std::generate_n`  
`std::swap_ranges`  
`std::reverse, std::reverse_copy`

## Removing & replacing elements

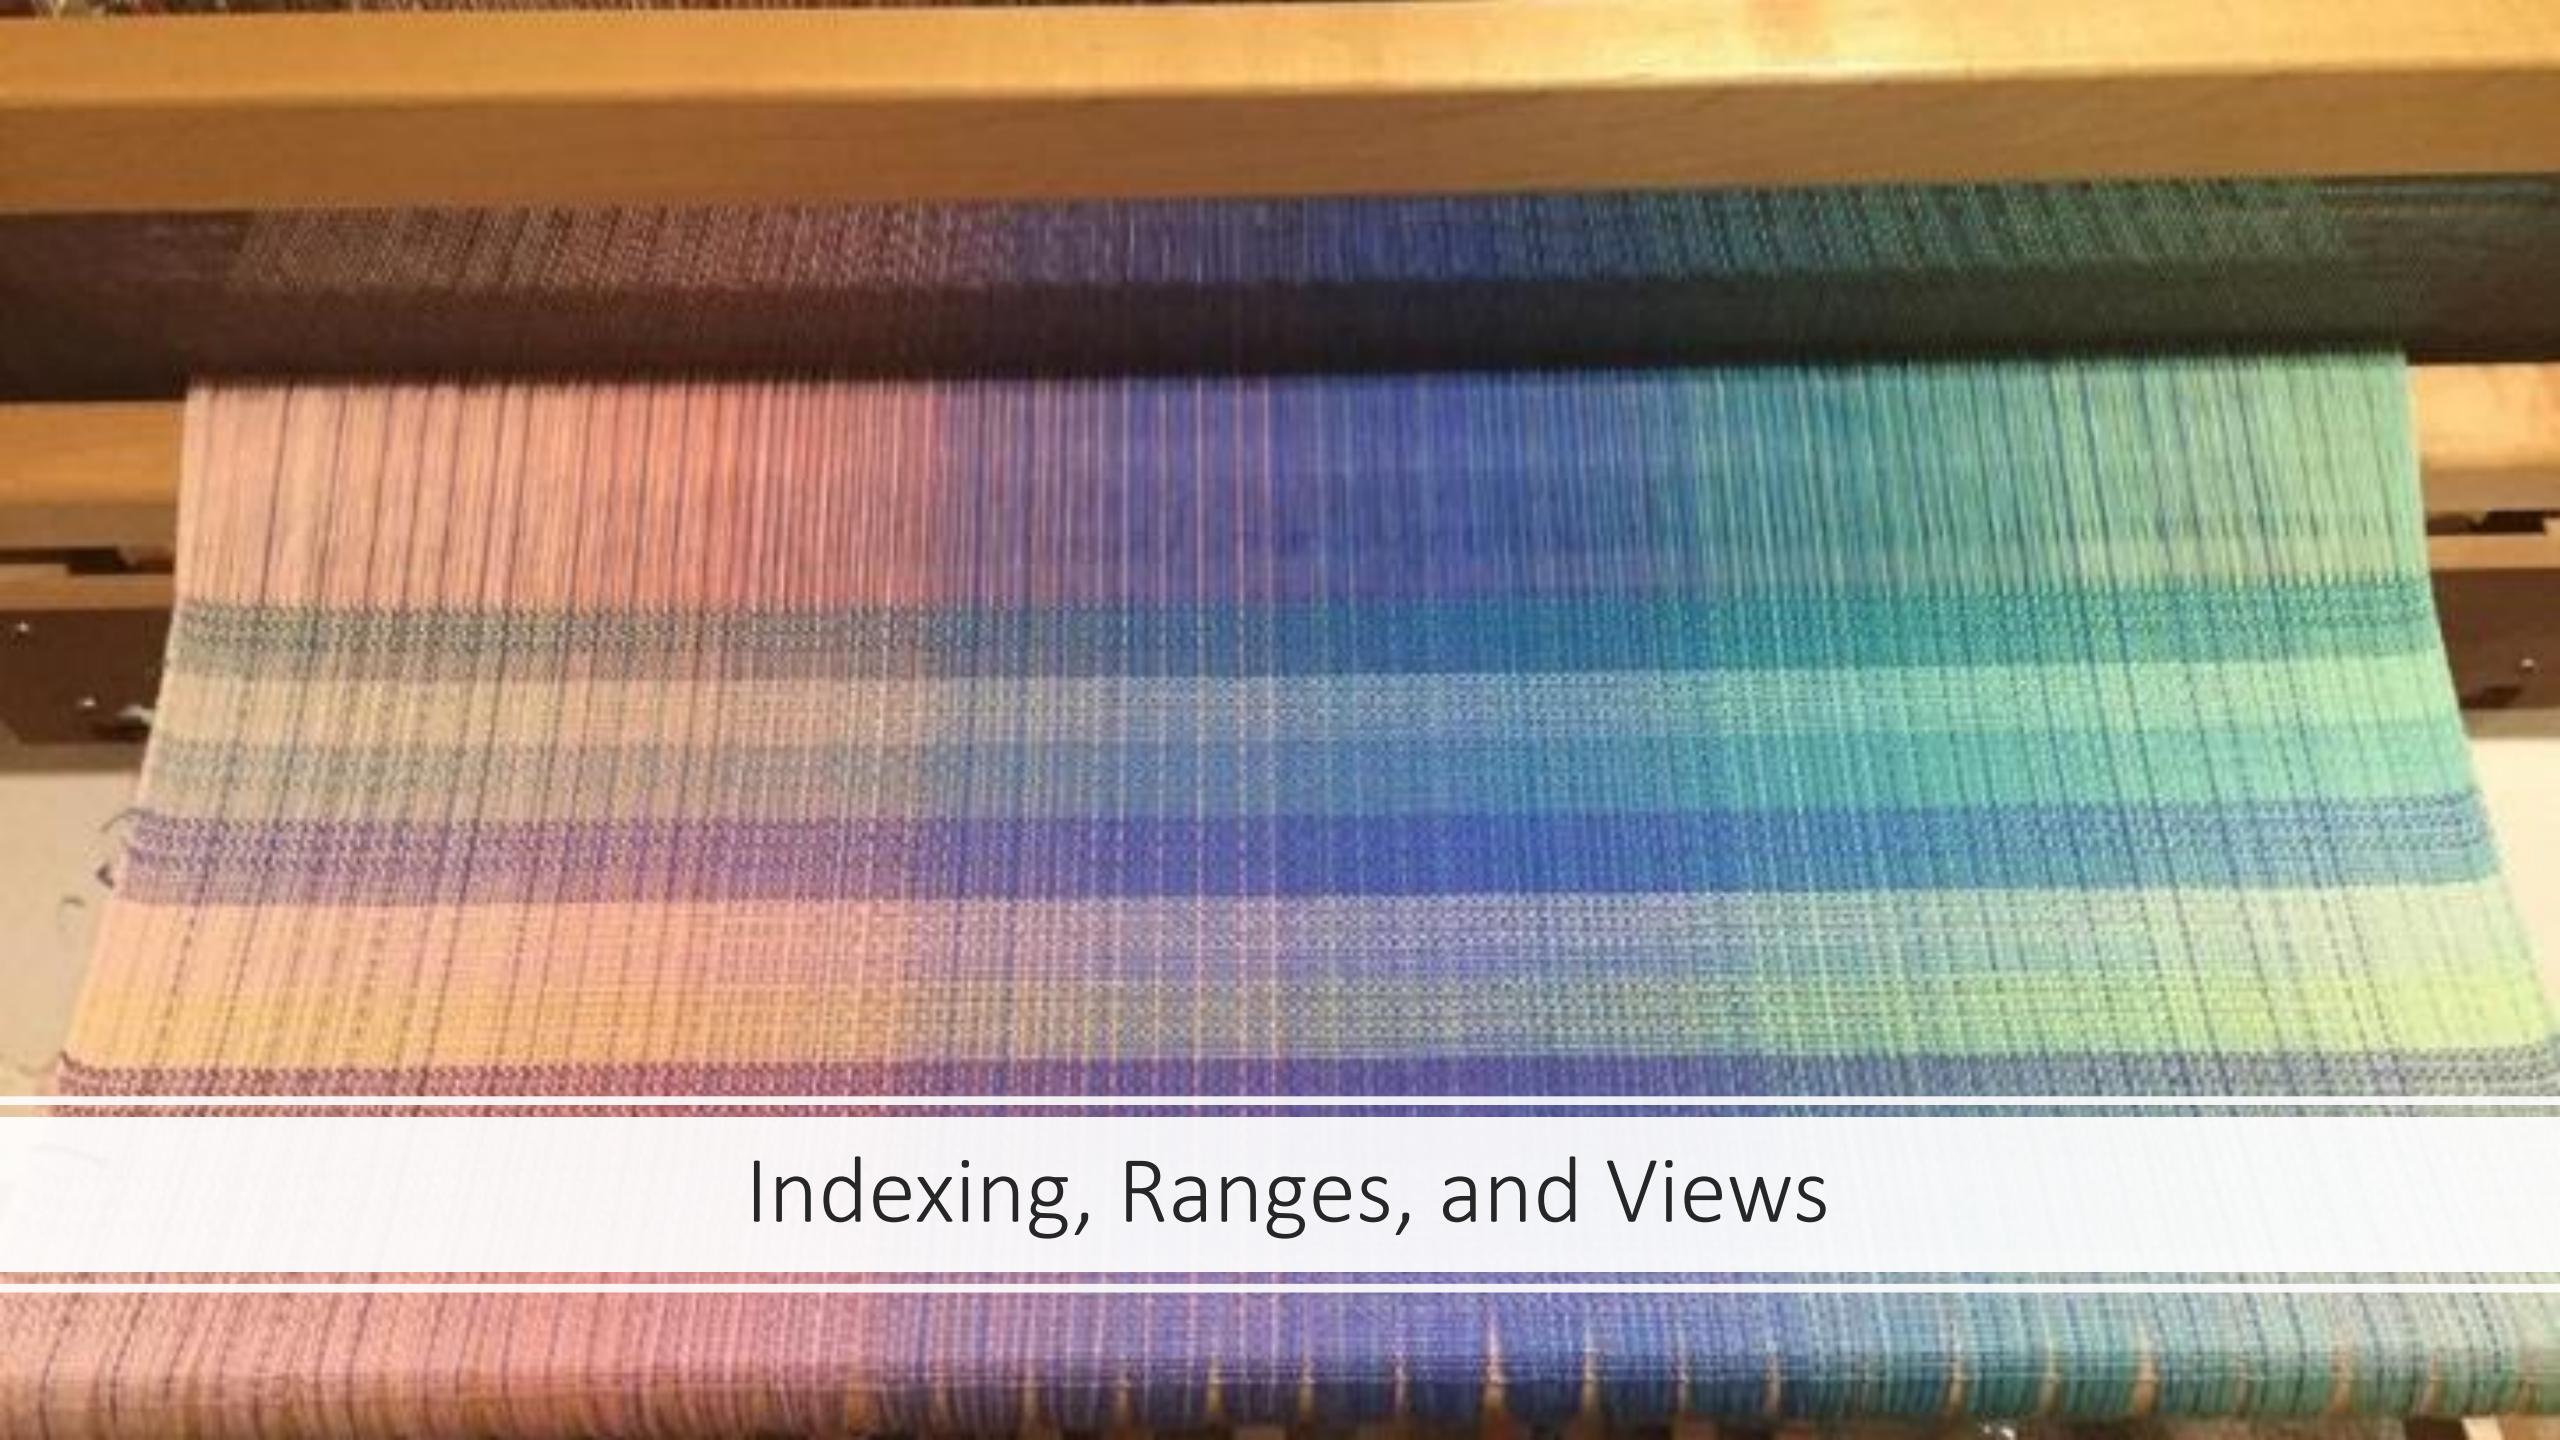
`std::remove, std::remove_if`  
`std::replace, std::replace_if, std::replace_copy, std::replace_copy_if`  
`std::unique / std::unique_copy`

## Reordering elements

`std::sort, std::stable_sort, std::partial_sort, std::partial_sort_copy`  
`std::rotate, std::rotate_copy, std::shift_left, std::shift_right`  
`std::partition, std::partition_copy, std::stable_partition`  
`std::nth_element`  
`std::merge, std::inplace_merge`

## Set operations

`std::includes, std::set_intersection, std::set_union`  
`std::set_difference, std::set_symmetric_difference`



# Indexing, Ranges, and Views

# How to find the index of an element?

With C++ `for` loops we have the index...

```
std::vector<double> v = {1, 2, 3, 4};  
for (int i = 0; i < 4 ; ++i) {  
    v[i] = f(i);  
}
```

...with parallel algorithms we do not...

```
std::transform(begin(v), end(v),  
             [](&const double& el) {  
                 return f(???);  
             });
```

# How to find the index of an element?

## Option 1: obtain index from address

With C++ `for` loops we have the index...

```
std::vector<double> v = {1, 2, 3, 4};  
for (int i = 0; i < 4 ; ++i) {  
    w[i] = f(i);  
}
```

...capture pointer to data by value (=) and compute the index from the memory address of the element...

```
std::transform(begin(v), end(v),  
    [v = v.data()](const double& el) {  
        ptrdiff_t i = &el - v;  
        return f(i);  
});
```

# How to find the index of an element?

## Option 2: use a counting iterator

A counting iterator is an iterator that wraps an index:

- [boost::counting\\_iterator](#)
- [thrust::counting\\_iterator](#)

... capture a pointer to the data by value (=) and use a counting iterator with the `std::for_each_n` algorithm...

```
thrust::counting_iterator<size_t> it{0};  
assert(*it == 0);  
++it;  
assert(*it == 1);  
  
std::for_each_n(it, v.size(),  
    [v = v.data()](size_t i) {  
        v[i] = f(i);  
});
```

# How to find the index of an element?

## Option 3: use C++20 Ranges and Views

The function *iota* from the C++20 collection of views defines an iterable sequence of numbers without actually allocating them.

Similarly, you can iterate over n-dimensional array indices using the view *cartesian\_product*.

```
auto ints = std::views::iota(0, 4);
std::for_each(par, begin(ints), end(ints),
    [v = v.data()](size_t i) {
        v[i] = f(i);
});

namespace stdv = std::views;
auto v = stdv::cartesian_product(
    stdv::iota(0, N), stdv::iota(0, M));

std::for_each(par, begin(v), end(v),
    [] (auto& e) {
        auto [i, j] = e;
    });
}
```

# How to find the index of an element?

## Summary

- Use pointer arithmetic (C++17)
  - In the lambda argument, pass the element by reference
  - Retrieve the index from the address of the element
- Use `counting_iterator` from a library (C++17)
  - Available in Thrust
  - Available in Boost
  - Available in other header files found on GitHub
- Use C++20 views
  - C++20 view `iota` for 1-D indexing. E.g. `views::iota(0, N).begin()`
  - C++23 view `cartesian_product` for n-D indexing.

# Background: C++20 Ranges and Views

A **Range** is an object that provides a pair of iterators denoting a range of elements, a `std::vector` is a range.

Sequential version of the C++ STL algorithms have versions that accept Ranges...

Parallel versions of the algorithms do **not!**

```
std::vector<double> v = {0, 1, 2, 3};  
auto b = v.begin();  
auto e = v.end();  
  
// Iterator version:  
std::transform(begin(v), end(v),  
 [](const double& el) { return 2. * el; });  
  
// Range version:  
std::ranges::transform(v,  
 [](const double& el) { return 2. * el; });
```

# Background: C++20 Ranges and Views

Views are lazy Range algorithms  
that produce elements as iterated  
over...

...we can use `views::iota` to  
generate a range of indices...

...that we can use with the `parallel`  
STL algorithms by using its  
iterators...

```
auto ints = std::views::iota(0, 4);
for (int i : ints) {
    v[i] = f(i);
}
```

```
std::for_each(par, begin(ints), end(ints),
              [v = v.data()](size_t i) {
                  v[i] = f(i);
});
```

# Background: C++20 Ranges and Views

Views algorithms compose via the "pipe" operator |...

```
auto seq = views::iota(0, N)
         | views::filter(is_prime)
         | views::stride(2)
         | views::transform(square);
for (auto i : seq) cout << i << ",";
// Prints: 4, 25, 121, ....
```

# C++20 and C++23 Views

- `views::all`
- `views::filter`
- `views::transform`
- `views::take`
- `views::join`
- `views::split`
- `views::zip`
- `views::counted`
- `views::reverse`
- `views::keys`
- `views::values`
- `views::cartesian_product`
- etc.

```
std::vector<int> xs{0, 1, 2, 3}, ys{4, 5, 6, 7};  
for (auto [x, y] : views::zip(xs, ys))  
    cout << "(" << x << "," << y << ")", ";  
// Prints: (0,4), (1,5), (2,6), (3,7)
```

# range-v3: <https://github.com/ericniebler/range-v3>

## Many Views and more for C++14 onwards

```
std::vector<int> w{4, 5, 6, 7};  
for (auto [x, y] : w | range::views::enumerate)  
    cout << "(" << x << ", " << y << "), ";  
  
// Prints: (0,4), (1,5), (2,6), (3,7)
```

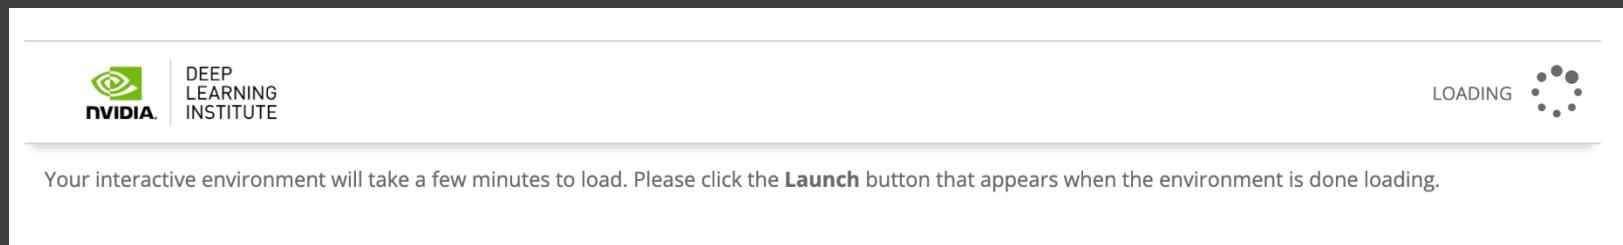
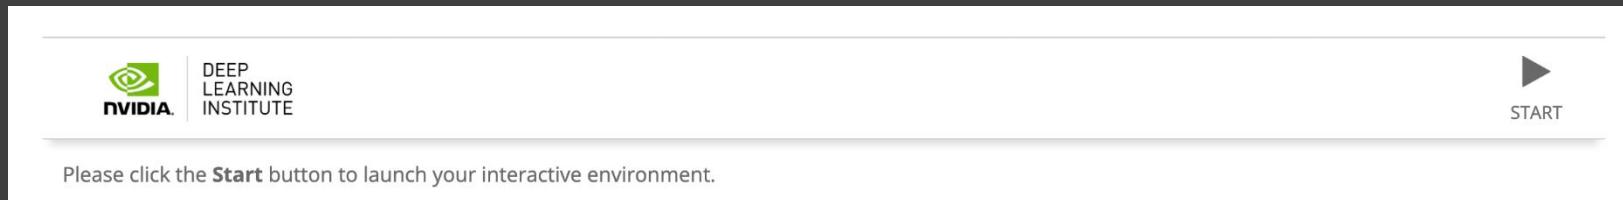
## References

- Tristan Brindle, An Overview of Standard Ranges, CppCon 2019
- Eric Niebler, [Ranges for the Standard Library](#), CppCon 2015

A close-up photograph of a loom's warp threads, which are woven into a vibrant, multi-colored fabric. The colors transition through a full spectrum, including red, orange, yellow, green, blue, and purple. The fabric has a distinct woven texture and is held in place by wooden beams at the top and bottom. The background is a plain, light-colored wall.

Interactive Materials

# Launch the Interactive Environment



# BLAS DAXPY: Double-precision AX + Y

Allocate memory...

```
std::vector<int> x(N), y(N);
```

Initialize x and y...

```
for (int i = 0; i < N; ++i) {  
    x[i] = ...;  
    y[i] = ...;  
}
```

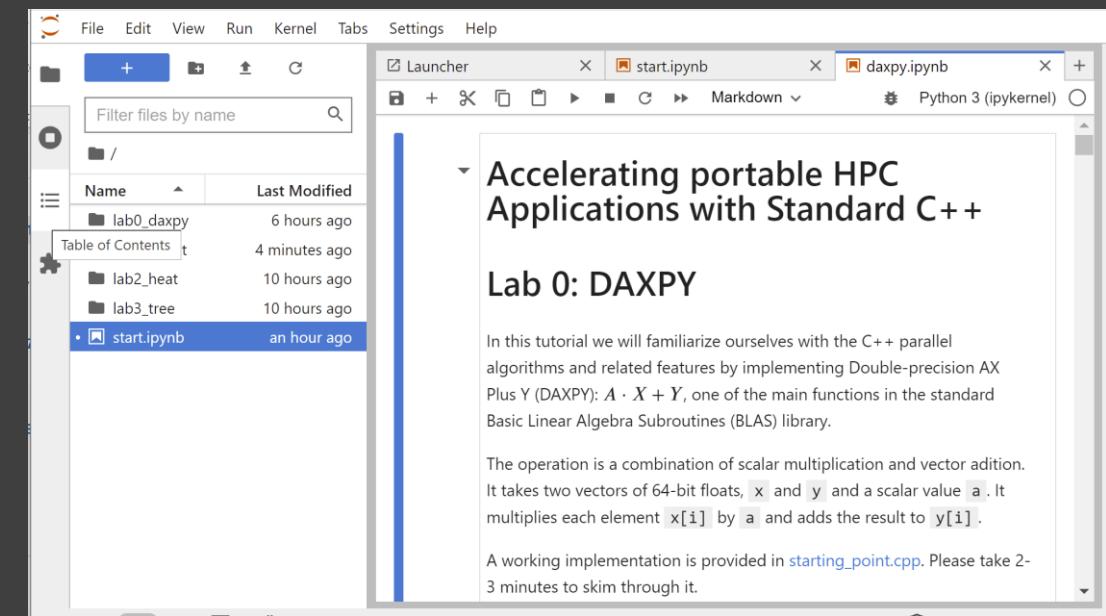
Update y...

```
for (int i = 0; i < N; ++i) {  
    y[i] += a * x[i];  
}
```

# BLAS DAXPY

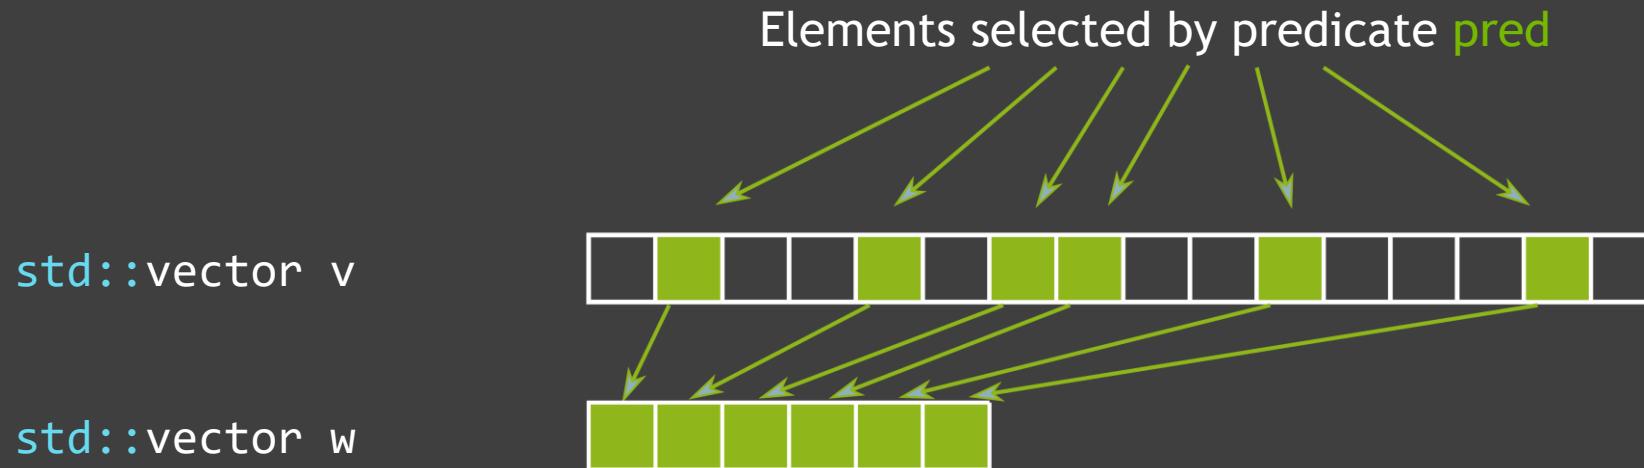
## 3 Exercises

- **Exercise 1:** rewrite the for-loop implementation of the BLAS DAXPY kernel using **sequential STL algorithms**
- **Exercise 2:** rewrite the for-loop implementation of the “initialization” kernel using **sequential STL algorithms** with any of the indexing approaches discussed in the tutorial
- **Exercise 3:** **parallelize** the STL versions of the application using the Execution Policies



# Extra Credit: Select

- **Exercise 1:** write the function `select`, which copies selected values from `v` to `w`.



- Sequentially no problem,  
but how to write a  
parallel algorithm ?

```
template<class UnaryPredicate>
std::vector<int> select( const std::vector<int>& v,
                        UnaryPredicate pred );
```

# Extra Credit: Select

- Step 1: Write out a binary-valued array for the values of the predicate.



# Extra Credit: Select

- Step 1: Write out a binary-valued array for the values of the predicate.
- Step 2: Compute the cumulative sum of this array.



# Extra Credit: Select

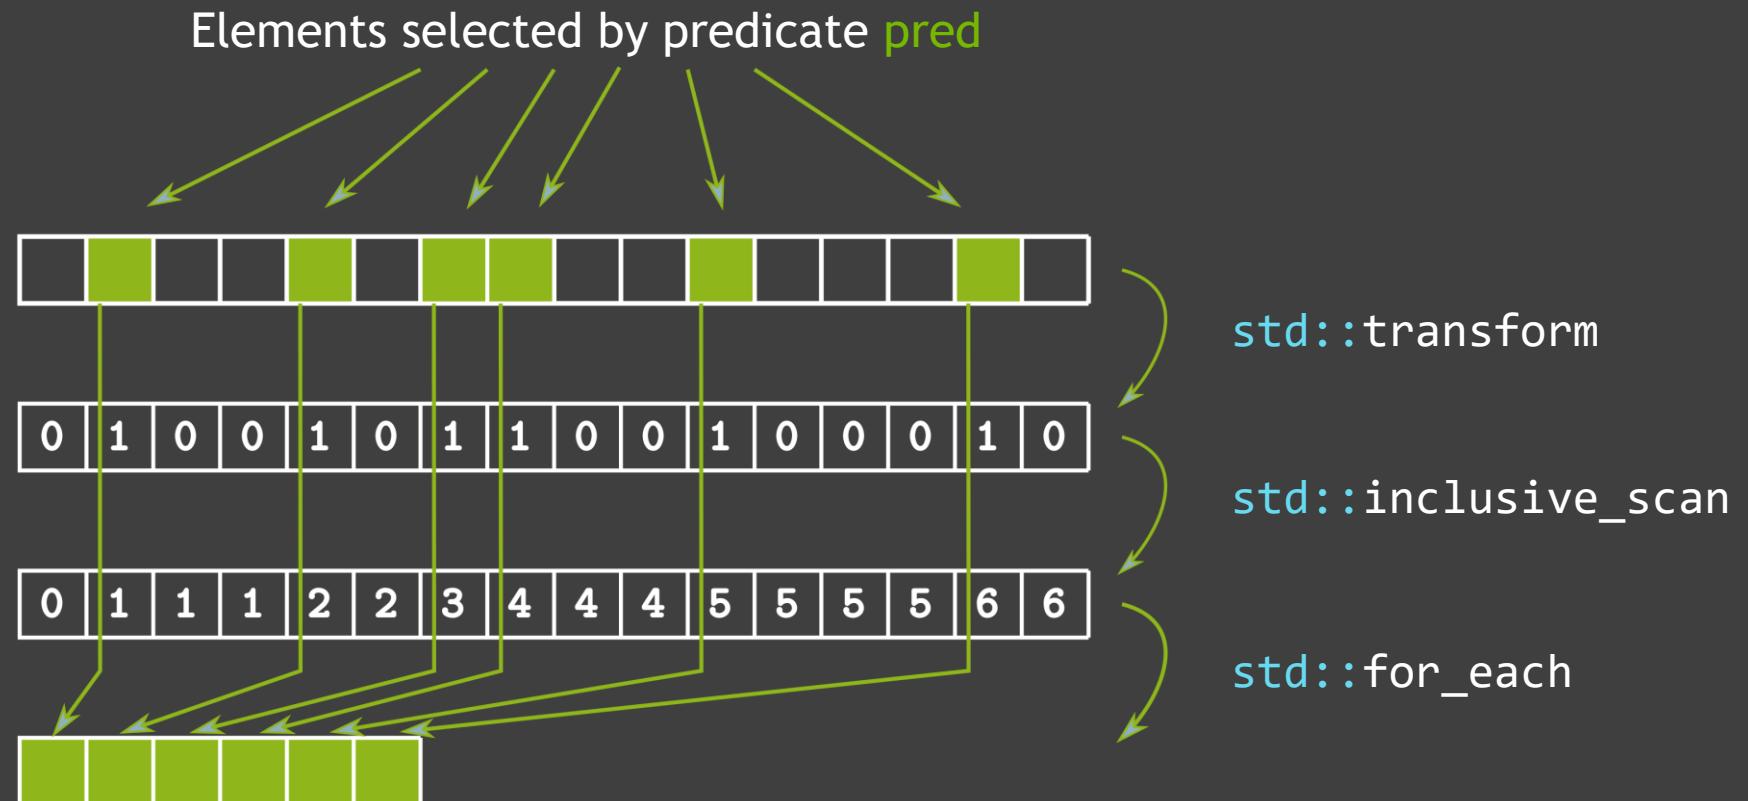
- **Step 1:** Write out a binary-valued array for the values of the predicate.
- **Step 2:** Compute the cumulative sum of this array.
- **Step 3:** Process vector v in parallel and use the cumulative sum to write to proper indices.

`std::vector v`

`std::vector w`

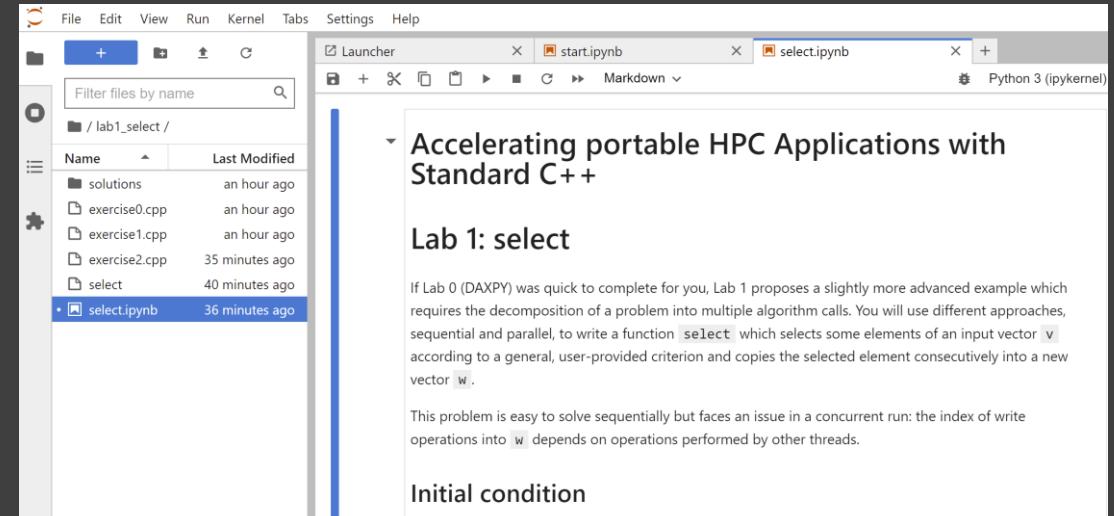
`std::vector index`

`std::vector w`



# Extra Credit: Select 3 Exercises

- **Exercise 1:** write a sequential version of the function `select` that uses the algorithm `std::copy_if` and a back inserter for vector `w`.
- **Exercise 2:** write a parallel version of the function `select` that works with two temporary vectors `v_sel` and `index`.
- **Exercise 3:** reduce the number of steps from 3 to 2 and avoid the creation of `v_sel` using the algorithm `transform_inclusive_scan`.



The screenshot shows a Jupyter Notebook interface with the following details:

- File Explorer:** On the left, it shows a directory structure for 'lab1\_select': 'solutions', 'exercise0.cpp', 'exercise1.cpp', 'exercise2.cpp', 'select', and 'select.ipynb'. 'select.ipynb' is currently selected.
- Launcher:** At the top, there are tabs for 'Launcher', 'start.ipynb', 'select.ipynb', and 'Python 3 (ipykernel)'.
- Notebook Content:** The main area displays the content of 'select.ipynb'.
  - Section Header:** `Accelerating portable HPC Applications with Standard C++`
  - Section Header:** `Lab 1: select`
  - Description:** A text block explaining that Lab 1 proposes a slightly more advanced example which requires the decomposition of a problem into multiple algorithm calls. It mentions using different approaches, sequential and parallel, to write a function `select` which selects some elements of an input vector `v` according to a general, user-provided criterion and copies the selected element consecutively into a new vector `w`. It notes that this problem is easy to solve sequentially but faces an issue in a concurrent run: the index of write operations into `w` depends on operations performed by other threads.
  - Text:** `Initial condition`

# The Interactive Materials



DEEP  
LEARNING  
INSTITUTE

Powered by:



This Lab 0 : 01 : 29 / 2 : 10 : 00

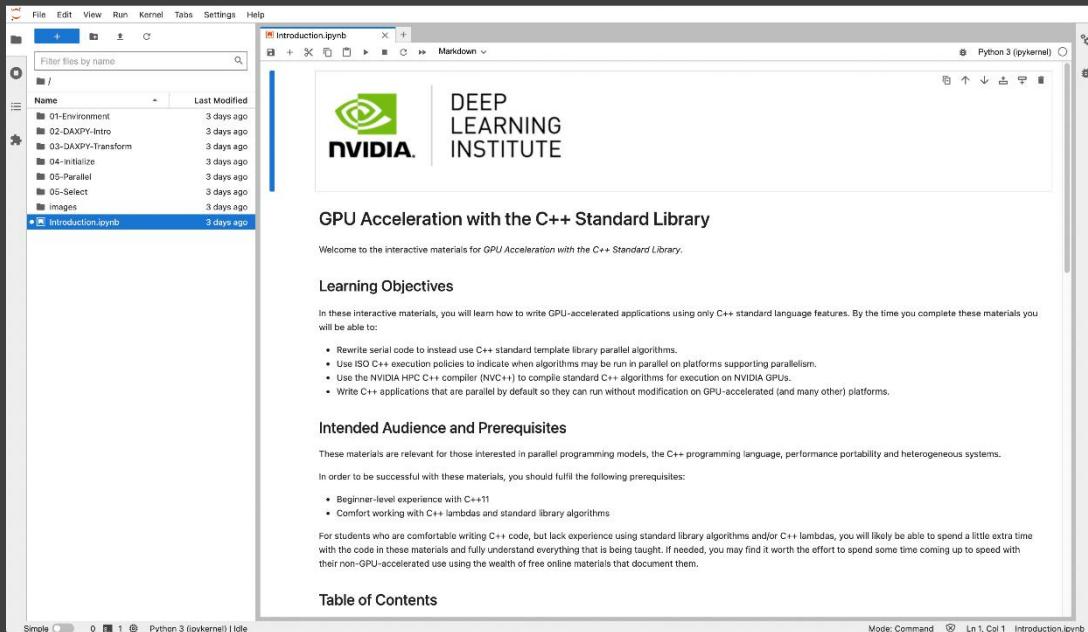


LAUNCH

STOP TASK

Your interactive environment will take a few minutes to load. Please click the **Launch** button that appears when the environment is done loading.

# The Interactive Materials



## Table of Contents

These materials consist of the following Jupyter notebooks, which should be worked through in order.

### Preliminaries

- [Introduction](#): This notebook, which introduces the interactive materials.
- [Your Compute Environment](#): Brief coverage of the component parts of the GPU-accelerated interactive environment you will be using for your work.

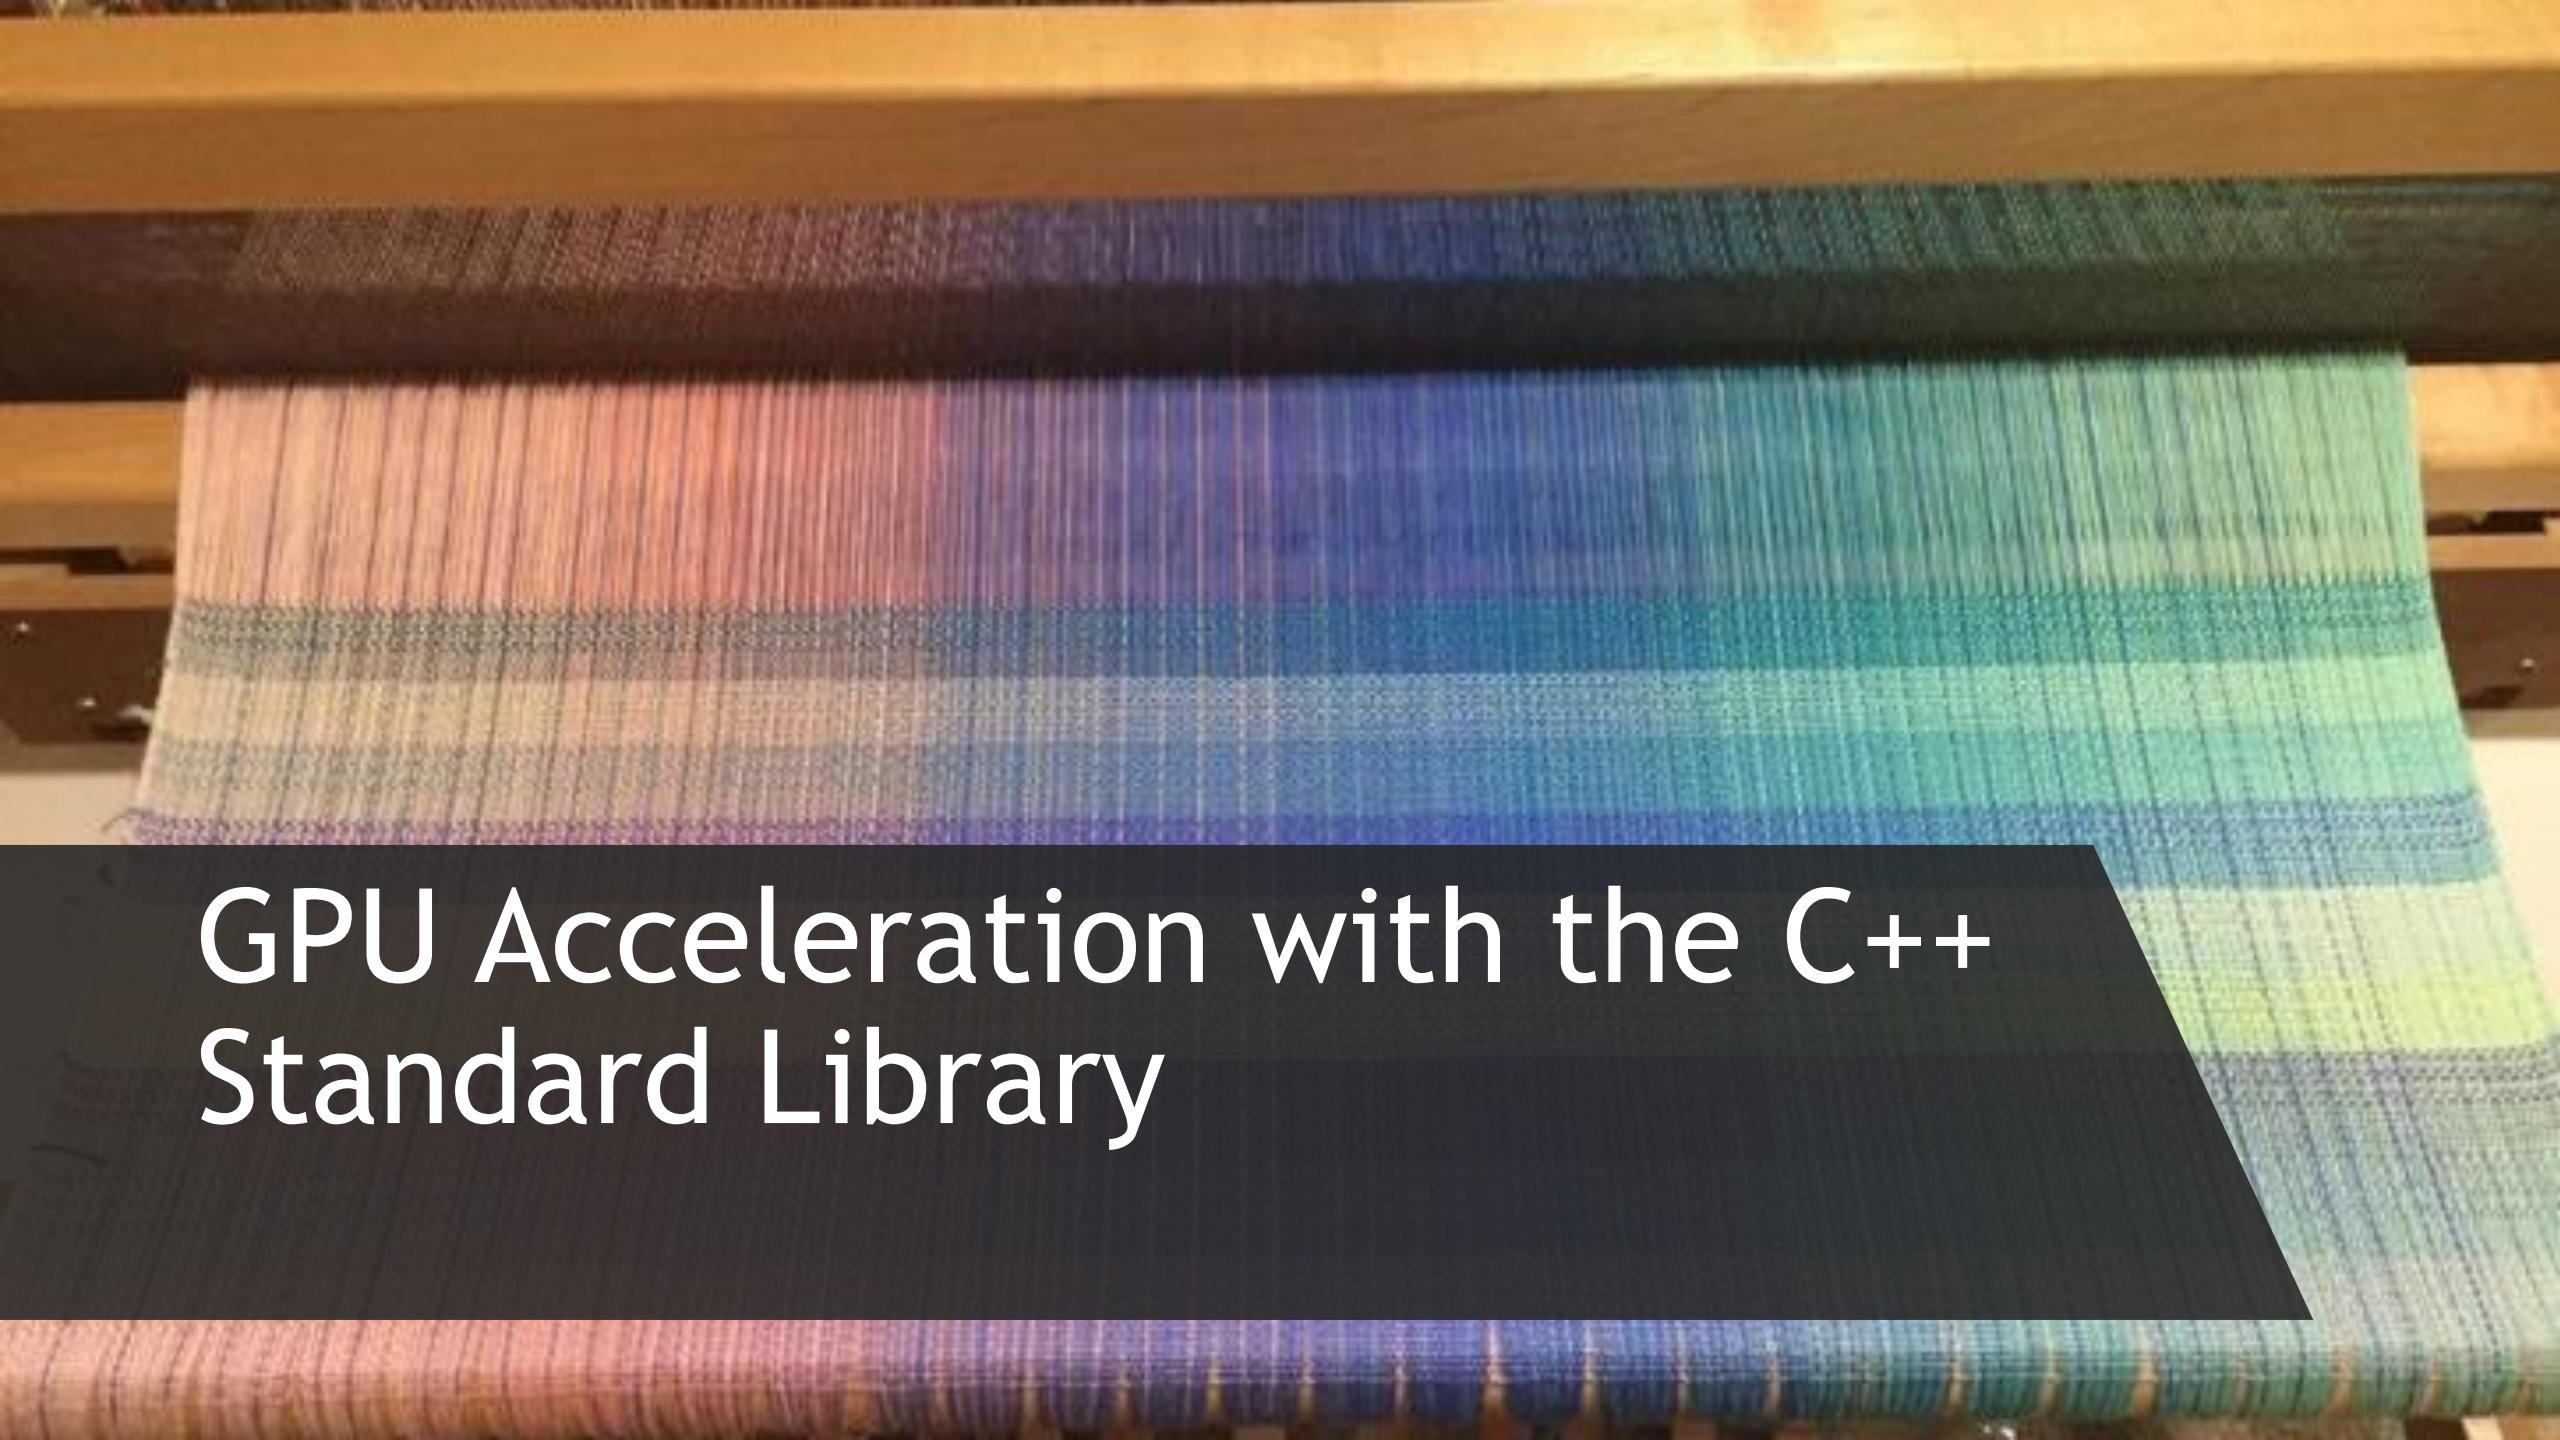
### Fundamentals of C++ Parallel Algorithms

- [Introducing DAXPY](#): The DAXPY operation, and an overview of a starting point sequential C++ DAXPY application.
- [From Raw DAXPY Loop to Serial C++ Transform Algorithm](#): The importance of the C++ algorithms library in "parallel first" code, and an exercise to use the `transform` algorithm.
- [From Raw Initialization Loop to Standard Library Algorithms](#): Refactoring `initialize` with the standard library to allow later for data creation on the GPU.
- [Parallelizing Daxpy and Initialization](#): Preparing the DAXPY program for parallel execution and observing massive performance gains on the GPU.

### Extra Credit: Solving Composite Problems with Algorithms

- [Select](#): Solving a composite problem with algorithms.





# GPU Acceleration with the C++ Standard Library